

## Assignment 2

### Q1: McCulloch-Pitts Neuron

In this question, I introduced the McCulloch-Pitts neuron, a simple mathematical model inspired by biological neurons. I drew the structure of a biological neuron and explained the correspondence:

- Dendrites correspond to input terminals.
- The cell body performs the summation of inputs.
- The axon transmits the output signal.
- Synapses represent the weights between neurons.

The McCulloch-Pitts neuron computes a weighted sum and applies a step function. I then showed how to build logic gates using this model:

- **AND gate:** weights = [1, 1], bias = -1.5
- **OR gate:** weights = [1, 1], bias = -0.5

Finally, I discussed why the **XOR gate cannot be implemented** with a single McCulloch-Pitts neuron — because XOR is not linearly separable.

---

### ■ Q2: MLP with Skip Connections (XOR Solver)

In this part, I analyzed a Multi-Layer Perceptron with a hidden neuron and a skip connection from input to output.

- I created a truth table showing that the network outputs correct XOR results.
- For **x<sub>3</sub>**, I derived the equation:  $x_3 = \text{step}(x_1 + x_2 - 1.5)$ . I drew the decision boundary: a line  $x_1 + x_2 = 1.5$ .
- For **x<sub>4</sub>**, the output is  $x_4 = \text{step}(x_1 - x_2 + 0.5)$ , so the decision line is  $x_1 - x_2 = -0.5$ .

These nonlinear boundaries together enable solving XOR.

---

### ■ Q3: Activation Functions

I introduced three popular activation functions:

- **Sigmoid:** outputs between 0 and 1
- **Tanh:** outputs between -1 and 1
- **ReLU:** outputs 0 for negative input, and x for positive input

I sketched all three on coordinate planes to show their behavior.

---

## ■ Q4: Shape of Parameters & Forward Pass

I described the matrix shapes of weights and biases in a standard feedforward neural network:

- **W1**: shape (d, h) — from input to hidden
- **b1**: shape (1, h)
- **W2**: shape (h, c) — from hidden to output
- **b2**: shape (1, c)

For input **x** (**1×d**), I wrote:

```
1 | f(x) = softmax(ReLU(xw1 + b1)w2 + b2)
```

And for batch input **X** (**n×d**):

```
1 | f(X) = softmax(ReLU(Xw1 + b1)w2 + b2)
```

## ■ Programming Part

In Python, I implemented 3 neural networks:

1. **ANN1**: 2 input → 2 output (no hidden layer), with softmax
2. **ANN2**: 2 input → 8 hidden (no activation) → 2 output, with softmax
3. **ANN3**: 2 input → 8 hidden (ReLU) → 2 output, with softmax

I used `matplotlib.pyplot.contourf()` to visualize decision boundaries.

- ANN1 has linear boundaries due to no hidden layer.
- ANN2 shows more complex patterns than ANN1, but still limited because it lacks non-linearity.
- ANN3 performs best, showing nonlinear decision boundaries thanks to the ReLU activation in the hidden layer.

I concluded that **multi-layer networks with non-linear activation** are superior because they can approximate complex functions and classify data that are not linearly separable.

## Assignment 4

### Q1 – Key Idea of Backpropagation

"The Backpropagation algorithm is a supervised learning technique used to train neural networks by minimizing the error between predicted and actual outputs.

It consists of a forward pass to compute the output, and a backward pass to compute the gradients using the chain rule.

These gradients are then used to update weights through gradient descent."

## Q2

"To answer Q2-2 efficiently, I wrote a Python script to simulate four steps of backpropagation training on a 1-2-1 MLP, using ReLU activation.

This helped automate the calculations of intermediate values such as activations, errors, gradients, and weight updates."

"Here's how the code works:"

### ◆ 1. Initialization

```
1 python复制编辑x = 0.7
2 d = 0.68
3 mu = 0.2
4 w1 = np.array([0.3, -0.3])
5 b1 = np.array([0.0, 0.0])
6 w2 = np.array([-0.1, 0.1])
7 b2 = 0.0
```

"The input `x` is 0.7, and the target output `d` is 0.68.

Weights and biases for the hidden layer (`w1`, `b1`) and output layer (`w2`, `b2`) are initialized manually.

The learning rate `mu` is set to 0.2."

#### 中文翻译:

"输入 `x` 是 0.7，目标值 `d` 是 0.68。

我手动初始化了隐藏层的权重和偏置 (`w1`, `b1`)，以及输出层 (`w2`, `b2`)。

学习率设为 0.2。"

### ◆ 2. Forward Pass and Backward Pass (Loop over 4 steps)

```
1 python复制编辑for t in range(steps):
2     z1 = w1 * x + b1
3     z = relu(z1)
4     z2 = np.dot(w2, z) + b2
5     y = z2
```

"In each training step, I compute the hidden layer activation `z` using ReLU, and the output `y` is computed as a linear combination of `z` and weights `w2`."

### ◆ 3. Compute Errors and Gradients

```
1 python复制编辑delta2 = y - d
2 delta1 = relu_derivative(z1) * w2 * delta2
3
4 deltaw2 = mu * delta2 * z
5 deltaw1 = mu * delta1 * x
```

"Then I compute the output error `delta2` and backpropagate it to get `delta1`.

These are used to compute weight updates `deltaw1`, `deltaw2`, and corresponding bias updates."

#### 中文翻译:

"我计算输出误差 `delta2`，再通过链式法则计算 `delta1`。

然后用这些误差计算每一层的权重和偏置更新量。"

### ◆ 4. Update Weights and Log Results

```
1 python复制编辑w1 -= deltaw1
2 w2 -= deltaw2
3 b1 -= mu * delta1
4 b2 -= mu * delta2
```

"After computing the gradients, I update the weights and biases.

I also log all intermediate results such as weights, activations, deltas, and gradients in a table format."

#### 中文翻译:

"完成梯度计算后，我更新了每一层的权重和偏置。

我还把所有中间结果（如权重、激活值、误差、梯度等）记录到表格中。"

### ◆ 5. Result Table Output

"The result is a clear, step-by-step table showing all variables for each training step, which matches the requirement of Q2-2.

This approach ensures accurate, repeatable calculations and saved time compared to manual derivation."

#### 中文翻译:

"最终结果是一个清晰的表格，记录了每一轮训练的所有变量，完全符合Q2-2题目的要求。

相比手动推导，这种方法更高效、更准确，也易于复查。"

"In the second part of the assignment, I implemented a neural network from scratch in Python using NumPy to solve the two-spiral classification task. I generated two classes of spiral data in 2D and used a fully connected network to separate them."

"The network has:

- 2 inputs
- Two hidden layers (each with 16 units)
- 1 output (sigmoid for binary classification)"

"I trained the model using the same BP principles I derived earlier."

### 中文翻译:

"作业的第二部分是用Python和NumPy从零实现神经网络，以解决双螺旋分类问题。我生成了两类二维螺旋数据，并使用一个前馈神经网络对其进行分类。"

"网络结构包括:

- 两个输入
- 两层隐藏层（每层16个神经元）
- 一个输出节点（sigmoid用于二分类）"

"我使用了和手推公式一致的反向传播方法进行训练。"

---

## ◆ Slide 6: Code Details and Results

"In the code:

- `train_nn()` runs forward and backward propagation for 2000 epochs
- `forward_pass()` is used to plot the decision boundary
- Loss was calculated using mean squared error"

"I plotted:

- A training loss curve (which shows convergence)
- A decision boundary plot (showing the learned separation)"

### 中文翻译:

"在代码中:

- `train_nn()` 函数执行前向和反向传播共2000轮
- `forward_pass()` 用于绘制决策边界
- 损失函数使用的是均方误差（MSE）"

"我绘制了:

- 训练损失曲线（显示模型收敛过程）
- 决策边界图（显示模型对螺旋数据的分离效果）"

## ◆ Slide 7: Reflection

"It was challenging to implement the BP algorithm fully by hand.  
I had to carefully calculate each derivative and verify the gradient shapes.  
Normalizing the data and tuning the learning rate were crucial for successful training."

### 中文翻译:

"手动实现完整的BP算法是有挑战性的。  
我需要非常仔细地计算每一层的导数并确保梯度维度一致。  
对数据进行归一化、调整学习率是训练成功的关键。"

## Assignment 5

### Q1 – Convolutional Neural Network Architecture

"The question provided a CNN architecture with two convolutional layers, one pooling layer, and two fully connected layers.  
I calculated the number of weights and biases for each layer based on kernel sizes, channels, and units."

#### ▪ Weights & Biases:

- Conv1:  $4 \times 4 \times 3 \times 8 = 384$  weights + 8 biases
  - Conv2:  $8 \times 8 \times 8 \times 16 = 8192$  weights + 16 biases
  - FC1:  $1024 \times 64 = 65536$  weights + 64 biases
  - FC2:  $64 \times 8 = 512$  weights + 8 biases
- Total: 74,624 weights, 96 biases**

#### ▪ Output Shapes:

- Input:  $84 \times 84 \times 3$
- Conv1:  $84 \times 84 \times 8$
- Conv2:  $26 \times 26 \times 16$  (due to stride 3 and valid padding)
- Pooling:  $8 \times 8 \times 16$
- FC1: 64, FC2: 8, Softmax: 8

#### ▪ Receptive Field:

- Final receptive field after pooling =  **$17 \times 17$**

## Q2 – 1D Convolution Output Derivation

"Given a 1D input signal and specified filters, I manually computed each output using valid padding, stride 1, and no activation."

### ▪ Conv1 Filters:

- Filter1 [0,1,0] → output: [2, 6, 7, 8, 9]
- Filter2 [0,0,1] → output: [6, 7, 8, 9, 1]
- Filter3 [1,0,0] → output: [1, 2, 6, 7, 8]
- Filter4 [1,0,1] → output: [7, 9, 14, 16, 9]

### ▪ Conv2:

- 1 filter, kernel size = 3×4, all weights = 1
- Result: [75, 99, 102]

## Q3 – Unrolling Deep RNNs

"For this question, I described how deep RNNs are unrolled through time. Each time step feeds into multiple RNN layers vertically. Horizontally, the hidden states propagate across time, and final outputs are taken from the last time step of the final layer."

### 编程

## 1. Importing Libraries

```
1 python复制编辑import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 import torchvision
5 import torchvision.transforms as transforms
```

"I start by importing PyTorch and torchvision libraries. These provide modules for defining neural networks, activation functions, and data transformations."

### 中文翻译:

"首先导入了 PyTorch 和 torchvision 库，它们提供了定义神经网络、激活函数、数据转换等必要的模块。"

## ◆ 2. Define the CNN Class

```
1 python复制编辑class MyCNN(nn.Module):
2     def __init__(self):
3         super(MyCNN, self).__init__()
4         ...
```

"I defined a class called `MyCNN`, which inherits from `nn.Module`.

Inside the constructor `__init__`, I define all the layers of my convolutional neural network."

### 中文翻译：

"我定义了一个类 `MyCNN`，继承自 PyTorch 的 `nn.Module`，  
在构造函数 `__init__()` 中定义了CNN模型的所有网络层。"

## ◆ 3. Network Layers

```
1 python复制编辑self.conv1 = nn.Conv2d(1, 16, kernel_size=5, padding=2)
2 self.bn1 = nn.BatchNorm2d(16)
3 self.pool1 = nn.MaxPool2d(2)
4
5 self.conv2 = nn.Conv2d(16, 32, kernel_size=5, padding=2)
6 self.bn2 = nn.BatchNorm2d(32)
7 self.pool2 = nn.MaxPool2d(2)
8
9 self.fc1 = nn.Linear(32*7*7, 128)
10 self.dropout = nn.Dropout(0.5)
11 self.fc2 = nn.Linear(128, 10)
```

"The CNN has:

- 2 convolutional layers (with batch normalization and max pooling)
- A fully connected layer with dropout
- A final output layer with 10 classes for MNIST"

"I used kernel size = 5, and padding = 2 to maintain image size."

### 中文翻译：

"模型包括：

- 两个卷积层（每个后接 BatchNorm 和 MaxPooling）
- 一个全连接层 + Dropout（防止过拟合）
- 最后一层为10个类别输出（用于MNIST）"



## ◆ 4. Forward Function

```
1 python复制编辑def forward(self, x):
2     x = self.pool1(F.relu(self.bn1(self.conv1(x))))
3     x = self.pool2(F.relu(self.bn2(self.conv2(x))))
4     x = x.view(-1, 32 * 7 * 7)
5     x = self.dropout(F.relu(self.fc1(x)))
6     x = self.fc2(x)
7     return x
```

"The forward method defines how input data flows through the network:

1. Conv → BN → ReLU → Pool
2. Conv → BN → ReLU → Pool
3. Flatten
4. Fully connected layer + dropout
5. Output layer"

### 中文翻译：

" `forward()` 函数定义了数据在网络中的流动方式：

1. 卷积 → BN → ReLU → 池化
2. 卷积 → BN → ReLU → 池化
3. 拉平为向量
4. 进入全连接层（带 Dropout）
5. 输出层给出分类结果"

## ◆ 5. Data Preprocessing

```
1 python复制编辑transform = transforms.Compose([
2     transforms.ToTensor(),
3     transforms.Normalize((0.5,), (0.5,))
4 ])
5 trainset = torchvision.datasets.MNIST(...)
6 trainloader = torch.utils.data.DataLoader(...)
```

"Here I downloaded the MNIST dataset and applied normalization using torchvision. Images are converted to tensors and pixel values are normalized to [-1, 1]."

### 中文翻译：

"使用 `torchvision` 下载MNIST数据集，并进行标准化处理：图像被转换为张量，并将像素值归一化到 [-1, 1] 区间。"

## ◆ 6. Training Setup

```
1 python复制编辑model = MyCNN()
2 criterion = nn.CrossEntropyLoss()
3 optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
```

"I used:

- `CrossEntropyLoss` as the loss function (suitable for multi-class classification)
- `Adam` optimizer with learning rate = 0.001"

中文翻译:

"使用交叉熵损失函数（适用于多分类任务）和Adam优化器，学习率设置为0.001。"

---

## ◆ 7. Training Loop

```
1 python复制编辑for epoch in range(num_epochs):
2     for i, (images, labels) in enumerate(trainloader):
3         outputs = model(images)
4         loss = criterion(outputs, labels)
5
6         optimizer.zero_grad()
7         loss.backward()
8         optimizer.step()
```

"The model is trained for several epochs.

In each iteration:

- Compute prediction
- Calculate loss
- Backpropagate and update weights"

中文翻译:

"模型进行了多轮迭代训练，每轮包括:

- 前向传播计算输出
  - 计算损失
  - 反向传播并更新权重"
-

## ◆ 8. Evaluation and Accuracy

```
1 python复制编辑with torch.no_grad():
2     correct = 0
3     total = 0
4     for images, labels in testloader:
5         outputs = model(images)
6         _, predicted = torch.max(outputs.data, 1)
7         correct += (predicted == labels).sum().item()
```

"After training, I evaluated the model accuracy on the test dataset.

Predicted labels are compared with true labels to compute the total accuracy."

### 中文翻译：

"训练完成后，我在测试集上评估模型准确率。

使用 `torch.max()` 获取最大概率对应的预测类别，并与真实标签进行比较。"